

# SQL: The Bad Cont'd

## Ordering Results:

- historically just an extra can be added to the top level query
- now: for some years, SQL had a LIMIT k clause, often used with ORDER BY (usually importance / relevance) "top k" (e.g. top 10 salary) due to non-determinism

## Null Values

Value inapplicable: bad schema design (compromise)

- better design exists, but poor performance

value unknown: can be replaced by any domain value satisfying integrity constraints

→ many possibilities (possible worlds)

how to answer queries?

→ true in all worlds  $W$  of an incomplete  $D$

Certain Answer

$$Q(D) = \bigcap_{W \text{ world of } D} Q(W)$$

→ answer common to all possible worlds

computationally infeasible in general!

SQL gives a crude approximation.

expressions: NULL as param to op makes result NULL  
 pred/comp. : 3-valued logic (crude approx. of V-Val.)  
 set ops: unique special value for duplicates  
 aggregate ops: doesn't count (value inapplicable)

Comps w/ NULL return UNKNOWN

$\wedge$	T	U	F	$\vee$	T	u	F	$\neg$	
T	T	u	F	T	T	T	T	T	F
u	u	u	F	u	T	u	u	u	u
F	F	F	F	F	T	u	F	F	T

$x \text{ null}$

$(x = 0) \text{ OR } (x \neq 0) = \text{UNKNOWN} \neq \text{TRUE}$   
<sub>unk                      unk</sub>

IS TRUE / IS FALSE / IS UNKNOWN

WHERE  $\langle \text{cond} \rangle$  is shorthand for WHERE  
 $\langle \text{cond} \rangle$  IS TRUE

select aid, name from author where url is NULL  
 count(URL) counts only non-NULL URLs  
 count(\*) counts rows

Outer Joins

Allows Null-padded answers that fail to satisfy a conjunct in a conjunction

FROM R  $\leftarrow_{j\text{-type}}$  JOIN S ON C <sup>FULL/LEFT/RIGHT/INNER</sup>

$\{ (x, y, z) \mid R(x, y) \wedge S(y, z) \vee (z = \text{NULL} \wedge R(x, y) \wedge \neg (\exists z. S(y, z))) \}$  (LEFT/FULL)

$\{ (x, y, z) \mid (x = \text{NULL} \wedge S(y, z) \wedge \neg (\exists x. R(x, y))) \}$  (RIGHT/FULL)

select aid, pub from author left join wrote on aid=author

1 ChTo98  
5 —

select aid, count(publication) as pubs  
from author left join wrote on aid=author  
#pub per author group by aid

NULLs are a necessary evil  
Can always be avoided, but possibly at cost of efficiency

can't escape NULLS in practice (schema evolution)

'connect to cs348'