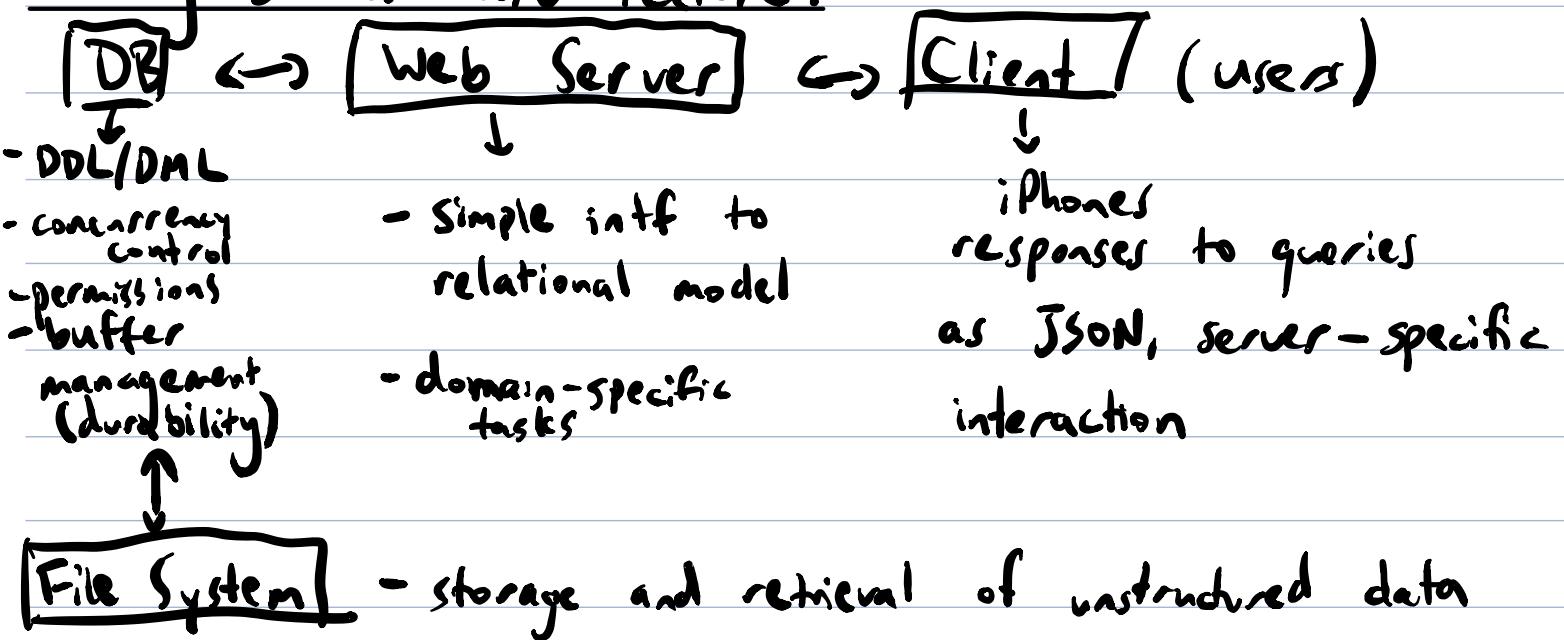


DB Tuning and Physical Design: Basics of Query Execution

We assume our tool uses a client-server architecture. Historically, db server code was implemented over a file system. In this course, assume a built-in ability to work with a 2-column table containing filenames and corresponding value of the file (sequence of bytes) allowing random access.

Usually 3-tier architecture:



Query

- ① Optimization / compilation
→ Output: Query plan
- ② Execution
→ Output: results (tables)

Straight to ② for query evaluation.

Goal: develop a simple relational calculator that answers range-restricted queries (bottom-up)

Considerations:

- physical storage (physical design)
- Compute answers to complex queries
- manage intermediate results

how to execute queries?

① parsing, type checking, etc

② SQL translated to relational algebra

③ optimization

⇒ generates an effective query plan

⇒ uses statistics collected about the stored data

④ plan execution

⇒ access methods to access stored relations

⇒ physical relational operators to combine relations

Need:

- query

- conceptual info (tables (schema))

- physical design (which data structs?)

- characteristics of result data (how big?)

⇒ statistically

Relⁿ algebra: $(U, R_0, \dots, R_{k_1}, \times, \sigma_p, \pi_v, \cup, -)$ {select } project

each R_i needs to be a name of a materialized view.

Projection

Definition

$$\pi_V(R) = \{ (x_{i_1}, \dots, x_{i_n}) \mid (x_1, \dots, x_n) \in R \wedge i_j \in V \}$$

where V is an ordered list of column numbers.

Selection

Definition

$$\sigma_\varphi(R) = \{ (x_1, \dots, x_n) \mid (x_1, \dots, x_n) \in R \wedge \varphi(x_1, \dots, x_n) \}$$

Product

e.g. Account \times Bank
(order matters!)

Can consider a simple VM implementation of the relⁿ algebra (nested loop).

Union: both sides need to have the same # of columns

$$\Pi_{1,2}(\sigma_{2=\text{CHK}}(\text{Account})) \cup \Pi_{1,2}(\sigma_{2=\text{SAV}}(\text{Account}))$$

Difference

e.g. $\Pi_{1,4}(\text{Account}) - \Pi_{1,4}(\sigma_{4,6}(\text{Account} \times \text{Bank}))$
Accounts without banks?

If we had bag semantics, we could add a unary operator for unique results only (duplicate elimination)

Theorem (Codd)

For every D.I. R.C. query there is an equivalent relational algebra expression.

$$\text{RCToRA}(R; (x_1, \dots, x_k)) = R;$$

$$\text{RCToRA}(Q \wedge x_i = x_j) = \sigma_{i=x_j}(\text{RCToRA}(Q))$$

$$\text{RCToRA}(\exists x_i \cdot Q) = \pi_{\text{FV}(Q) - \{x_i\}}(\text{RCToRA}(Q))$$

...

How do we (mostly) avoid storing intermediate results?

- ⇒ cursor open/fetch/close interface
- every implementation of an R.A. operator
- i) produces answers this way

2) gets answers from children this way
We make at least one physical implementation per operator.

This DB implementation will be very slow!

- 1) use disk-based data structs for efficient searching (indexing)
- 2) use better algs to implement operators (SORTING, HASHING)
- 3) rewrite expression to equivalent, more efficient one

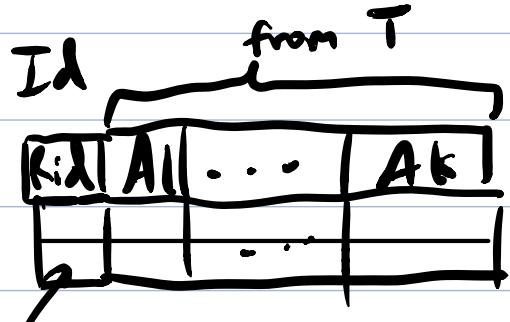
Atomic Relations and Indexing

- when an index $R_{\text{index}}(x)$ (where x is the search attribute) is a variable, we replace $\sigma_{x \in c}(R)$ with accessing $R_{\text{index}}(x)$ directly
- otherwise, check all file blocks holding tuples for R
- even if $\text{id } X$ available, searching entire tables may be faster if:
 - relation is small, or
 - rel^n is large, but we expect most of the tuples to satisfy the condition

Standard physical design

create table T ...

- instantly creates materialized view with an internally generated identifier
- create view Id as (select * from T) materialized



address of record encoding that tuple in the file

- this is automatically done
- vendor may provide commands to manipulate this physical design
(create index ...)

create index I (A_{i_1}, \dots, A_{i_s}) on T [clustered]

clustering: tuples with similar value of attribute are stored together in the same block

e.g. create index Id () on T clustered

create index I1 (A_3, A_4) on T

non-clustering (secondary): creates materialized view with columns (A_3, A_4, rid) where rid is a foreign key to the clustered index