


deduplication can be implemented in various ways:


- hashing
- sorting

in any case, it's expensive (need temp data struct)



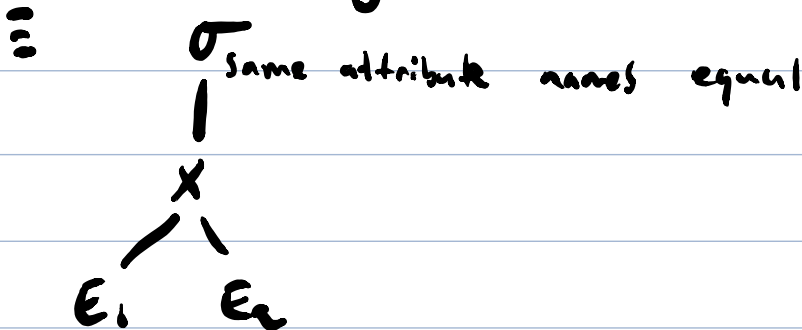
Additional Operators

- Deduplicate(E)
- $R(\#i = c)$
- $E_1 \bowtie E_2 \equiv$



- $E_1 \bowtie E_2$

↑ natural join

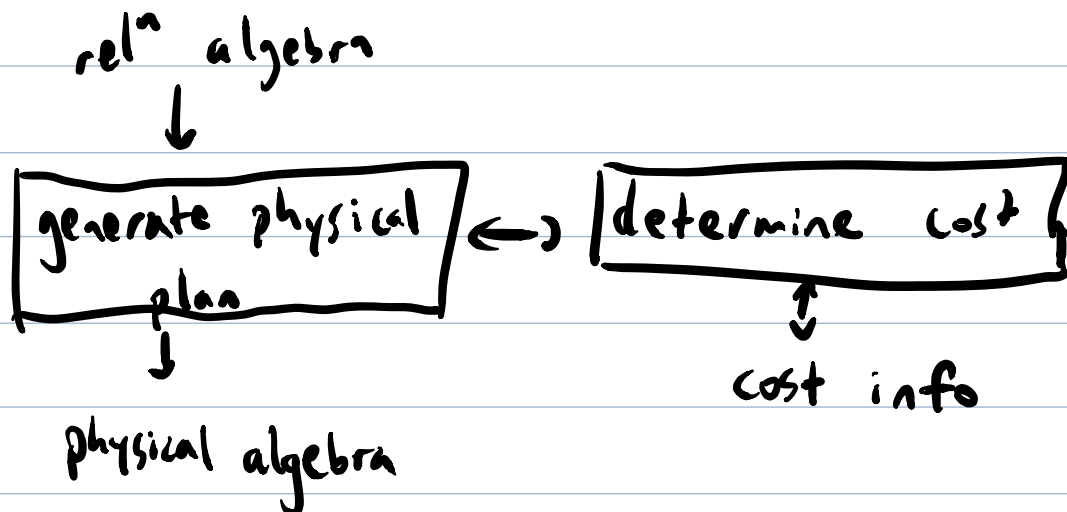


- $E_1 \div E_2$ (like universal quantification)

syntactic sugar only


Finding an optimal query plan is computationally infeasible \Rightarrow we look for a reasonable one

- db engine has statistical metadata, helps estimate relative cost of plans



finding all equivalent plans is undecidable in general.

- expensive even for conjunctive queries
 - \hookrightarrow the join-ordering problem: given a selection of conditions



find best ordering of the child joins (NP-hard)

\Rightarrow use heuristics

simple cost model for disk I/O, assume:

uniformity: all possible attrib equally likely

independence: likelihood of attrib. having a value in a tuple does not depend on values of other attribute

⇒ pretty draconian assumptions

$b(R)$:= blocking factor for stored relⁿ R
(describes # of entries in a page on disk or in B+ - tree)

π_1

σ_1

Course Ind


selecting N tuples from relⁿ R using primary index has a cost of $2 + \frac{N}{b(R)}$

in this case $2 + \frac{100}{b(\text{mark})} = 2 + \frac{100}{50} = \boxed{4}$


$\frac{|\text{Mark}|}{100} = 100$ tuples for PHYS

π_1

σ_1



$$\text{Cost} = 2 + N = \underline{21}$$



$$\text{Cost}(\sigma_c(E)) = (1 + \epsilon_c) \text{Cost}(E)$$

$$\text{Cost}(R \bowtie S) = \text{Cost}(R) + \frac{|R|}{b} \text{Cost}(S)$$

$$\approx |R| |S|$$

$$\text{Cost}(R \bowtie S) = \text{Cost}(R) + d_s |R| \quad \text{index join}$$

Sort-merge join

$$\text{Cost}(R \bowtie S) = \text{Cost}(\text{Sort}(R)) + \text{Cost}(\text{Sort}(S))$$

$$\text{where } \text{Cost}(\text{Sort}(E)) = \text{Cost}(E) + \frac{|E|}{b} \lg \frac{|E|}{b}$$

Definition (Selectivity)

$$\text{sel}(\sigma_c(R)) = \frac{|\sigma_c(R)|}{|R|}$$

Always good: push selections down

$$\sigma_p(E_1 \bowtie E_2) = \sigma_p(E_1) \bowtie \sigma_p(E_2)$$

push projections down
replace products by joins