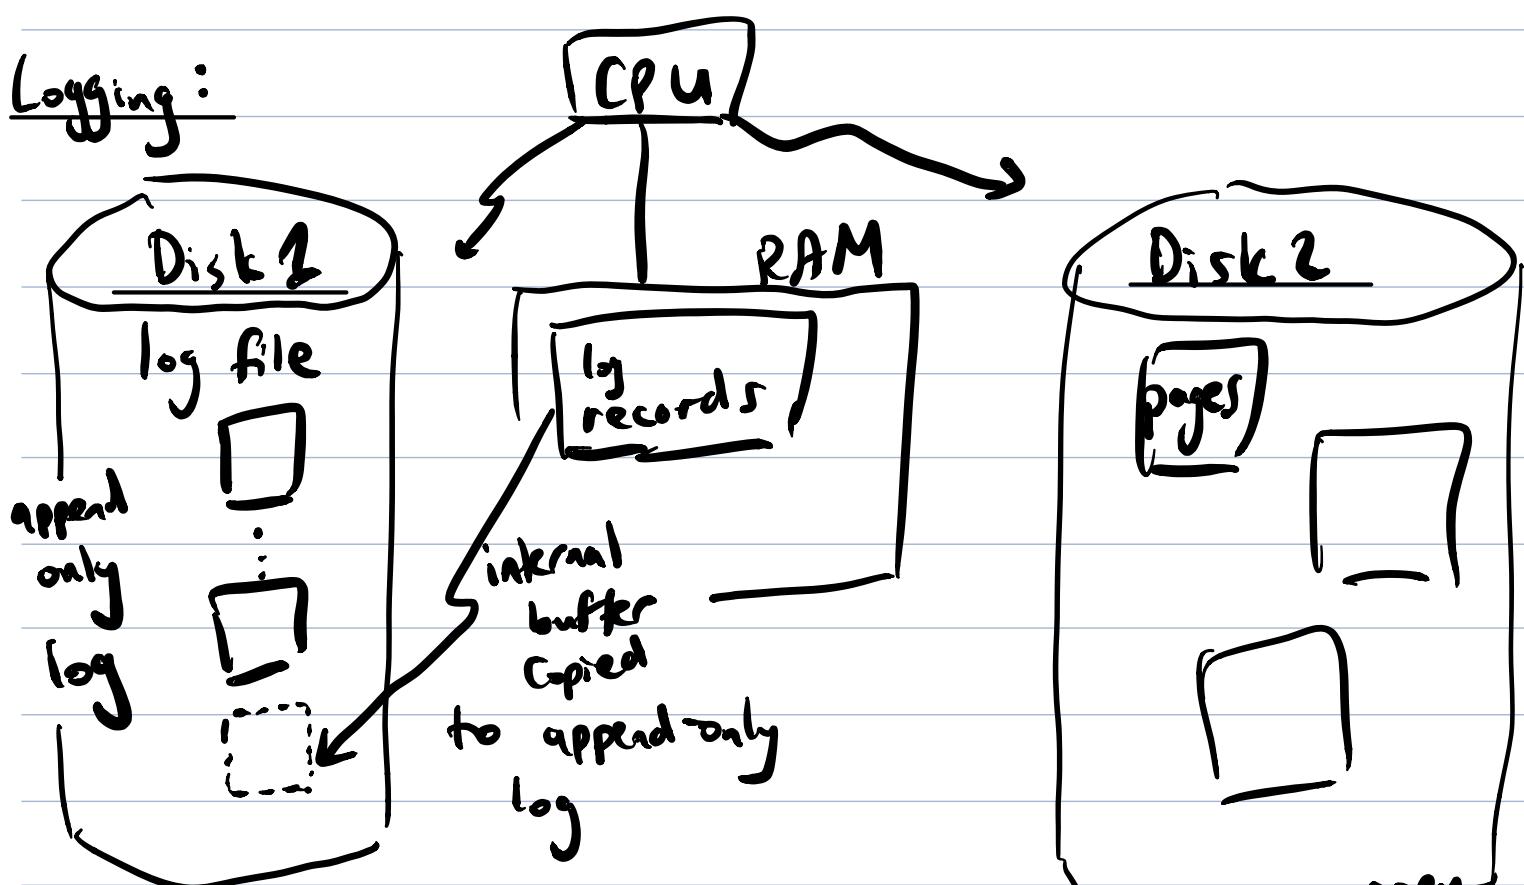


Shadowing : not updating primary copy of database
 end transaction \Rightarrow assign root of changes to
 new version

Logging:



log head \rightarrow T_0, begin
 (oldest part) $T_0, X, 99, 100$ } log
 \vdots

log tail
 (newest part) T_3, commit

recovery algorithm reconciles durable primary
 copy w/ the log

Need to follow write-ahead logging (WAL):

1) undo rule: log record for update is
 written to log disk before corresponding

data (page) is written to the main disk
(guarantees atomicity)

2) redo rule: all log records for a tx
are written to log disk before commit
(guarantees durability)

recover from lost RAM:

- 1) undo everything in WAL (\uparrow)
- 2) redo everything (\downarrow)

may not need to undo everything:

- undo aborted tx
- abort/active tx
- redo committed tx

how do we prevent WAL from growing
unbounded?

A: Checkpointing

- periodically enter mode: intention to write a checkpoint (not normal processing)
 - no new transactions
 - wait for active tx to finish
 - Wait for dirty pages to be written back to primary disk
 - THEN write a checkpoint log record and flush to log disk

→ now we know at that checkpoint record, db is in a good state and we don't need to look further up

Next big optimization: Checkpointing is a transaction, that records info allowing dbms to infer how far further back to look

Summary:

- ACID guarantees correctness on concurrent db access and data storage
- C/I based on serializability
 - (Schedulers, transaction manager)
- D/A based on recovery manager
 - ⇒ Synchronous writing is too inefficient
 - ⇒ replaced by synchronous writes to a Log and WAL

tuning: adjust physical/conceptual schemas to adjust to changed workloads or requirements on performance characteristics

- everything changes!
- ⇒ never outside of "maintenance modes"
- change naturally happens to underlying workload

Good design and tuning requires understanding the database workload.

Definition

A workload description contains:

- most important {queries, updates} and their frequency
- desired performance goal for each query or update

db tuning goal: make a set of applications execute as fast as possible

- optimize for response time?
- optimize for overall throughput?
(e.g. #tx / second)

as DBA, how can we affect performance?

- make queries faster (data structs, clustering, replication)
- make updates faster (locality)
- minimize congestion due to concurrency

We talked about standard physical design:

- every pk: primary index

- zero or more secondary indices

Storage options: unsorted (heap) file, sorted file, hashed file, dynamic hashing

Add indices: B+ trees, R trees, hash tables, ISAM, VSAM

IBM products (GOS)
B+ trees w/o splitting/merging

Multidimensional indexing
e.g. Efficiently find all cities in long/lat query latitude
multiple range querying

ISAM, VSAM are static indexing approaches:

- these make sense for read-heavy workloads
- like arrays
- column-store is a static indexing scheme with an array, at its core
- OLAP workloads

key takeaway: B-trees aren't the end of the story

Select * from emp where LN = 'Smith'

- if no indices for LN \Rightarrow all blocks of heap file must be scanned

Create index LNIndex on Emp (LN) [CLUSTER]

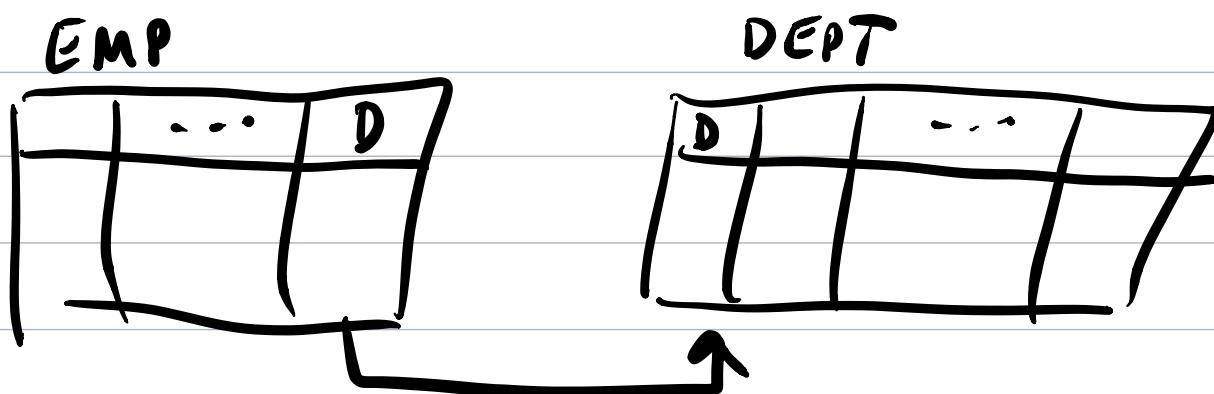
①

drop index LNIndex

②

- ①: - ↓ execution time for lastname selections
- ↑ execution time for insertions
- ↑ or ↓ exec. time for updates or deletions of tuples from emp
- ↑ space to represent emp

Co-clustering:



- interleave tuples from two relations in the same file
- useful for storing hierarchical data ($1:N$ relationships)

- effects on performance
 - 1) speed up joins, particularly FK joins
 - 2) sequential scans of either relations becomes



slower

join indexes: allow replacing joins by index lookups

materialized views: allow replace queries by index lookups

multi-attribute indices: complex search/join cond "x"

Problem 1: optimizer needs to know if/where to use such indices/views

Problem 2: balance cost of rematerialization and savings for queries.